

kaviyadevi 20106064

```
In [1]: #to import Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: #to import dataset
data1=pd.read_csv(r"C:\Users\user\Downloads\15_Horse Racing Results.CSV - 15_Hors
data1
```

```
Out[2]:
```

| | Dato | Track | Race Number | Distance | Surface | Prize money | Starting position | Jockey | Jockey weight | Countr |
|-------|------------|------------|----------------|----------|---------|----------------|----------------------|------------------|------------------|--------------|
| 0 | 03.09.2017 | Sha Tin | 10 | 1400 | Gress | 1310000 | 6 | K C Leung | 52 | Sverig |
| 1 | 16.09.2017 | Sha Tin | 10 | 1400 | Gress | 1310000 | 14 | C Y Ho | 52 | Sverig |
| 2 | 14.10.2017 | Sha Tin | 10 | 1400 | Gress | 1310000 | 8 | C Y Ho | 52 | Sverig |
| 3 | 11.11.2017 | Sha Tin | 9 | 1600 | Gress | 1310000 | 13 | Brett Prebble | 54 | Sverig |
| 4 | 26.11.2017 | Sha Tin | 9 | 1600 | Gress | 1310000 | 9 | C Y Ho | 52 | Sverig |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 27003 | 14.06.2020 | Sha Tin | 11 | 1200 | Gress | 1450000 | 6 | A Hamelin | 59 | Australi |
| 27004 | 21.06.2020 | Sha Tin | 2 | 1200 | Gress | 967000 | 7 | K C Leung | 57 | Australi |
| 27005 | 21.06.2020 | Sha Tin | 4 | 1200 | Gress | 967000 | 6 | Blake Shinn | 57 | Australi |
| 27006 | 21.06.2020 | Sha Tin | 5 | 1200 | Gress | 967000 | 14 | Joao Moreira | 57 | Ne Zealan |
| 27007 | 21.06.2020 | Sha Tin | 11 | 1200 | Gress | 1450000 | 7 | C Schofield | 55 | Ne Zealan |

27008 rows × 21 columns



In [3]:

```
#to display top 5 rows
data=data1.head(100)
data
```

Out[3]:

| | Dato | Track | Race Number | Distance | Surface | Prize money | Starting position | Jockey | Jockey weight | Country |
|-----|------------|---------|-------------|----------|---------|-------------|-------------------|-------------------------|---------------|---------------|
| 0 | 03.09.2017 | Sha Tin | 10 | 1400 | Gress | 1310000 | 6 | K C Leung | 52 | Sverige |
| 1 | 16.09.2017 | Sha Tin | 10 | 1400 | Gress | 1310000 | 14 | C Y Ho | 52 | Sverige |
| 2 | 14.10.2017 | Sha Tin | 10 | 1400 | Gress | 1310000 | 8 | C Y Ho | 52 | Sverige |
| 3 | 11.11.2017 | Sha Tin | 9 | 1600 | Gress | 1310000 | 13 | Brett Prebble | 54 | Sverige |
| 4 | 26.11.2017 | Sha Tin | 9 | 1600 | Gress | 1310000 | 9 | C Y Ho | 52 | Sverige |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 95 | 10.12.2017 | Sha Tin | 5 | 1200 | Gress | 18500000 | 13 | Francois-Xavier Bertras | 57 | Great Britain |
| 96 | 10.12.2017 | Sha Tin | 7 | 1600 | Gress | 23000000 | 11 | Ryan Moore | 57 | USA |
| 97 | 01.10.2017 | Sha Tin | 7 | 1000 | Gress | 3000000 | 10 | Brett Prebble | 59 | New Zealand |
| 98 | 22.10.2017 | Sha Tin | 7 | 1200 | Gress | 4000000 | 9 | Brett Prebble | 59 | New Zealand |
| 99 | 19.11.2017 | Sha Tin | 7 | 1200 | Gress | 4000000 | 3 | Brett Prebble | 56 | New Zealand |

100 rows × 21 columns



DATA CLEANING AND PREPROCESSING

```
In [4]: #
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Dato                   100 non-null   object
1   Track                  100 non-null   object
2   Race Number            100 non-null   int64
3   Distance                100 non-null   int64
4   Surface                 100 non-null   object
5   Prize money             100 non-null   int64
6   Starting position      100 non-null   int64
7   Jockey                  100 non-null   object
8   Jockey weight           100 non-null   int64
9   Country                 100 non-null   object
10  Horse age               100 non-null   int64
11  TrainerName             100 non-null   object
12  Race time               100 non-null   object
13  Path                    100 non-null   int64
14  Final place             100 non-null   int64
15  FGGrating               100 non-null   int64
16  Odds                    100 non-null   object
17  RaceType                100 non-null   object
18  HorseId                 100 non-null   int64
19  JockeyId                 100 non-null   int64
20  TrainerID               100 non-null   int64
dtypes: int64(12), object(9)
memory usage: 16.5+ KB
```

```
In [5]: #to display summary of statistics(here to know min max value)
data.describe()
```

```
Out[5]:
```

| | Race Number | Distance | Prize money | Starting position | Jockey weight | Horse age | Path | Final place |
|--------------|----------------|-------------|--------------|----------------------|------------------|--------------|------------|----------------|
| count | 100.000000 | 100.000000 | 1.000000e+02 | 100.000000 | 100.000000 | 100.000000 | 100.000000 | 100.000000 |
| mean | 6.910000 | 1446.000000 | 3.562200e+06 | 6.170000 | 55.870000 | 6.580000 | 1.510000 | 6.170000 |
| std | 2.099038 | 334.820923 | 4.486259e+06 | 3.440857 | 2.942736 | 1.35721 | 1.573101 | 3.440857 |
| min | 1.000000 | 1000.000000 | 9.200000e+05 | 1.000000 | 49.000000 | 3.000000 | 0.000000 | 1.000000 |
| 25% | 6.000000 | 1200.000000 | 1.380000e+06 | 3.000000 | 54.000000 | 6.000000 | 0.000000 | 4.000000 |
| 50% | 7.000000 | 1400.000000 | 1.950000e+06 | 6.000000 | 56.000000 | 7.000000 | 1.000000 | 6.000000 |
| 75% | 8.000000 | 1650.000000 | 3.000000e+06 | 9.000000 | 58.000000 | 8.000000 | 3.000000 | 9.000000 |
| max | 10.000000 | 2400.000000 | 2.300000e+07 | 14.000000 | 60.000000 | 9.000000 | 6.000000 | 12.000000 |

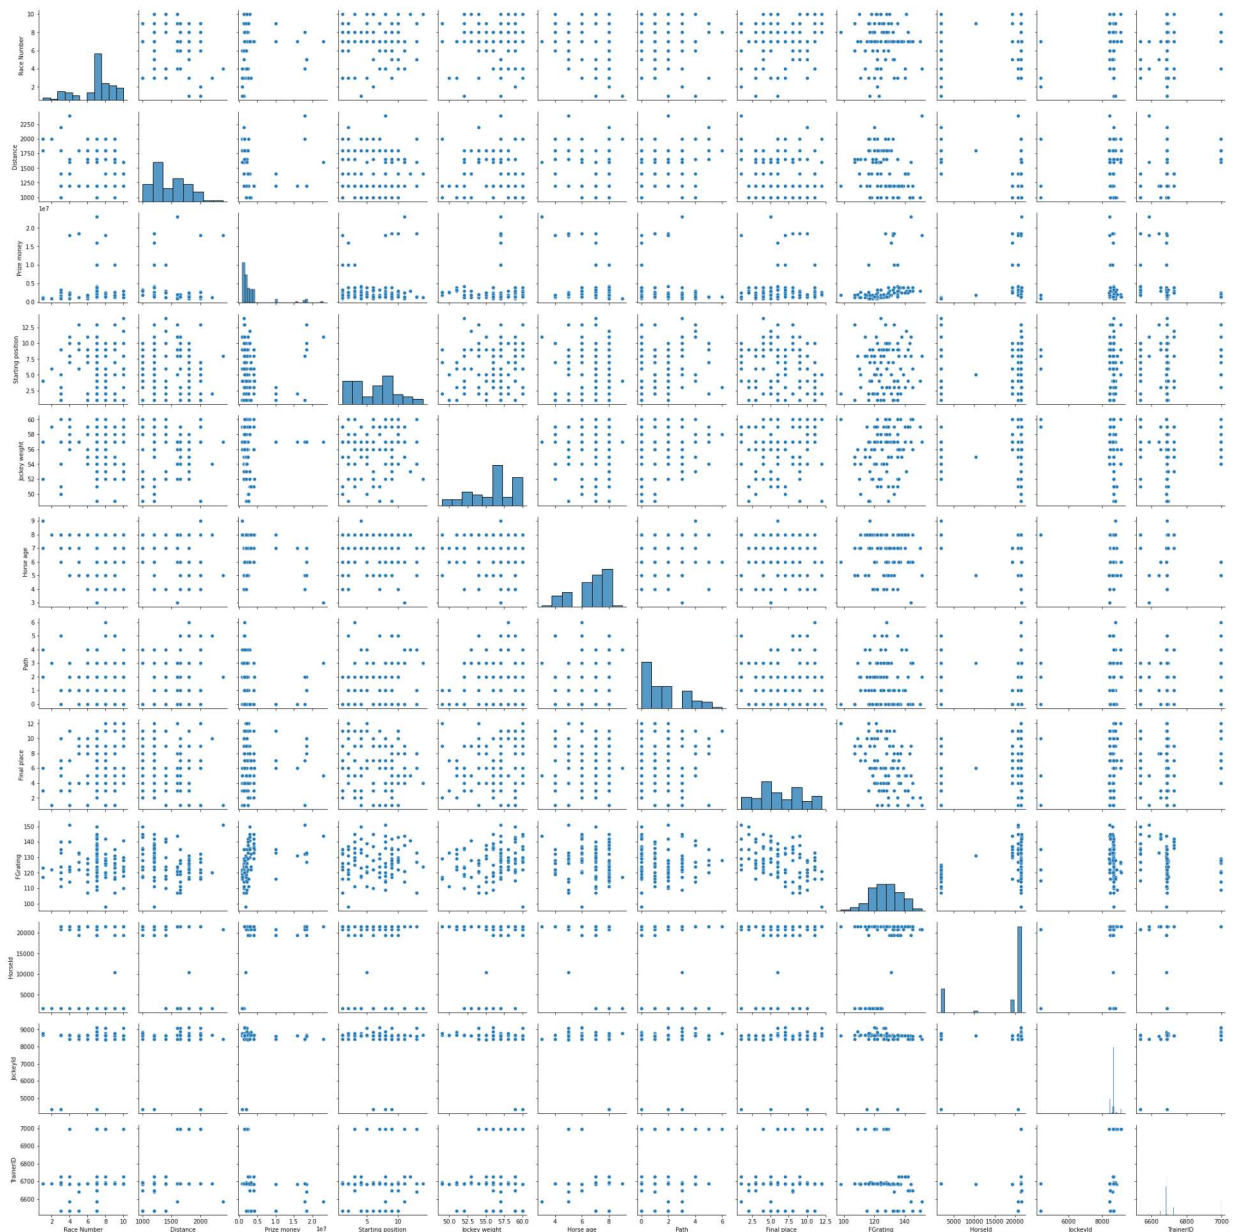
```
In [6]: #to display the column heading
data.columns
```

```
Out[6]: Index(['Dato', 'Track', 'Race Number', 'Distance', 'Surface', 'Prize money',
              'Starting position', 'Jockey', 'Jockey weight', 'Country', 'Horse age',
              'TrainerName', 'Race time', 'Path', 'Final place', 'FGrating', 'Odds',
              'RaceType', 'HorseId', 'JockeyId', 'TrainerID'],
              dtype='object')
```

EDA and DATA VISUALIZATION

```
In [7]: sns.pairplot(data)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x21c23c75910>
```




```
In [8]: sns.distplot(data['Final place'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```


```
Out[8]: <AxesSubplot:xlabel='Final place', ylabel='Density'>
```



```
In [9]: df=data[['Dato', 'Track', 'Race Number', 'Distance', 'Surface', 'Prize money',  
                'Starting position', 'Jockey', 'Jockey weight', 'Horse age', 'Path', 'Final place', 'FG rating', 'Horse sex', 'JockeyID', 'TrainerID']]
```

```
In [10]: sns.heatmap(df.corr())
```

```
Out[10]: <AxesSubplot: >
```



TRAINING MODEL

```
In [11]: x=df[['Prize money', 'Jockey weight', 'Horse age', 'Path', 'HorseId', 'JockeyId']
y=df['Final place']
```

```
In [12]: #to split my dataset into training and test

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [13]: from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[13]: LinearRegression()

```
In [14]: #to find intercept
print(lr.intercept_)
```

9.36303809762785


```
In [15]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[15]:

| | Co-efficient |
|----------------------|---------------|
| Prize money | 7.609949e-10 |
| Jockey weight | 7.739040e-02 |
| Horse age | 1.099069e-01 |
| Path | 2.010332e-01 |
| Horseld | 7.163400e-05 |
| Jockeyld | 6.222035e-04 |
| TrainerID | -2.204513e-03 |

```
In [16]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x21c2d7b8700>



```
In [17]: print(lr.score(x_test,y_test))
```

-0.03502949396008659

RIDGE AND LASSO REGRESSION

```
In [18]: from sklearn.linear_model import Ridge,Lasso
```

```
In [19]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[19]: Ridge(alpha=10)

```
In [20]: rr.score(x_test,y_test)
```

Out[20]: -0.036968758402129875

```
In [21]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[21]: Lasso(alpha=10)

```
In [22]: la.score(x_test,y_test)
```

```
Out[22]: -0.05353449065594229
```

```
In [24]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[24]: ElasticNet()
```

```
In [25]: print(en.coef_)
```

```
[-6.39536183e-10  1.42371537e-02  0.00000000e+00  0.00000000e+00
 6.02573468e-05  5.76873207e-04 -2.10922701e-03]
```

```
In [26]: print(en.predict(x_test))
```

```
[3.22351042 6.76993259 6.4345862  6.77236712 6.30965858 7.06892676
 5.6151163  6.229081  6.93259798 6.79868718 6.52042482 6.74670244
 6.78545053 6.63099513 6.71261943 6.90185187 6.17821072 6.73481609
 6.21962001 6.91299057 6.74506634 6.33037967 6.04416663 6.71261943
 6.76993259 4.69214  6.7788909  6.77742651 6.74173856 6.81821291]
```

```
In [27]: print(en.score(x_test,y_test))
```

```
-0.07732682931719137
```

```
In [28]: from sklearn import metrics
```

```
In [29]: print("Mean Absolute error",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolute error 2.445588877527731
```

```
In [30]: print("Mean Squared error",metrics.mean_squared_error(y_test,prediction))
```

```
Mean Squared error 9.06340826877716
```

```
In [31]: print("Root Mean Absolute error",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
Root Mean Absolute error 3.010549496151352
```

```
In [ ]:
```